

MASTER'S THESIS 2023

Drones for Sea Rescue: A Quality of Experience Study of Camera Control

Alexander Sandström



DEPARTMENT OF ELECTRICAL AND INFORMATION
TECHNOLOGY

LTH | LUND UNIVERSITY



EXAMENSARBETE
Elektro- och Informationsteknik

**Drones for Sea Rescue: A Quality of
Experience Study of Camera Control**

Drönare inom Sjöräddning: en Studie i
Operatörsupplevelse av Kamerakontroll

Alexander Sandström

Drones for Sea Rescue: A Quality of Experience Study of Camera Control

Alexander Sandström

alexander.h.sandstrom@gmail.com

May 9, 2023

Master's thesis work carried out at the Swedish Sea Rescue Society and the Department of Electrical and Information Technology, Lund University.

Supervisors: Fredrik Falkman, fredrik.falkman@ssrs.se
William Tärneberg, william.tarneberg@eit.lth.se

Examiner: Maria Kihl, maria.kihl@eit.lth.se

Abstract

Your abstract should capture, in English, the whole thesis with focus on the problem and solution in 150 words. It should be placed on a separate right-hand page, with an additional *1cm* margin on both left and right. Avoid acronyms, footnotes, and references in the abstract if possible.

Leave a *2cm* vertical space after the abstract and provide a few keywords relevant for your report. Use five to six words, of which at most two should be from the title.

Keywords: MSc, BSc, template, report, style, structure

Acknowledgements

If you want to thank people, do it here, on a separate right-hand page. Both the U.S. *acknowledgments* and the British *acknowledgements* spellings are acceptable.

Contents

1	Introduction	7
1.1	Drone use	7
1.2	SSRS Needs	7
1.3	Research Motivation	8
1.4	Scope	8
2	Background	9
2.1	Swedish Sea Rescue Society	9
2.1.1	Innovation	9
2.1.2	Drone Project: Eyes On Scene	9
2.2	Quality of Experience	10
2.3	Previous Work	10
3	Test Beds	11
3.1	Hardware	11
3.1.1	Fixed-Wing Drone	11
3.1.2	Gimbal Package	13
3.1.3	Hoverboard robot	13
3.1.4	Apriltag	13
3.2	Software	14
3.2.1	Pre-existing Software	14
3.2.2	Gimbal Control Interface	15
3.2.3	Result Script	16
3.2.4	Auto-follow mode	16
4	Experiments	19
4.1	QoE Lab Experiment	19
4.1.1	Task	19
4.1.2	Experimental Setup	19
4.1.3	Experimental Procedure	21

4.1.4	Evaluation	23
4.2	Known Sources of Error	25
4.2.1	Ecological validity	25
4.2.2	Frame detection accuracy	25
4.2.3	Difficulty of the task	26
4.3	Field testing	26
4.3.1	Goals	27
4.3.2	Preparations	27
4.3.3	Procedure	27
5	Results	29
5.1	Quantative measurements	29
5.2	Forms and interviews	31
6	Discussion	33
6.1	Video streaming	33
6.2	Native vs. web	33
6.3	Sources of error	34
6.3.1	Frame detection accuracy and interpolation	34
6.4	Experiment execution	34
7	Conclusion	35
8	Future Work	37
References		39
Appendix A	About This Document	43
A.1	Page Size and Margins	44
A.2	Typeface and Font Sizes	44
A.2.1	Headers and Footers	44
A.2.2	Chapters, Sections, Paragraphs	45
A.2.3	Tables	45
A.2.4	Figures	46
A.3	Mathematical Formulae and Equations	46
A.4	References	47
A.5	Colours	47
Appendix B	Language	49
B.1	Style Elements	49
Appendix C	Structure	51

Chapter 1

Introduction

This chapter will give a short introduction to use cases of UAVs and the work of SSRS along with the motivation and scope for the thesis work.

1.1 Drone use

With the recent advancements in technology, the use of drones has skyrocketed. Most notably they have been used in warfare for surveillance and attack purposes, but in recent years the civilian applications have also increased (SOURCES). Common civilian use cases are in cinematography, agriculture and delivery, and there is also a growing interest in using drones for search and rescue (GET MORE EXAMPLES).

Historically drones have been limited by the maximum distance it can travel from its operator, but with advancements in mobile network technology the use case for drones has grown stronger as they now have a much larger range and can be controlled from anywhere with an internet connection.

1.2 SSRS Needs

The SSRS has been conducting a research project called Eyes-On-Scene where drones are to be used to give better information to a rescue crew on their way to the scene of the accident. As part of the project a mission control software in which one can manage a fleet of drones and see their video feeds has been implemented. The drone camera is controlled by a so called region of interest (ROI) point, which is a GPS-location where the camera aims automatically. As the plane wiggles a lot during flight, stabilization is a necessity. With the camera pointed at an ROI, the plane's sensors can compensate for its movements resulting in a much more stable image. The act of uploading a new ROI requires a point to be selected on a map, which then can be uploaded to the drone. This forces a context switch for the operator from the

video feed to the map. To make it easier to operate the gimbal, SSRS want to be able to adjust the gimbal angle with direct controls, such as with a joystick or the arrow keys. The SSRS also wants an estimate of the latency between the drone and the operator, which they currently don't know. To illustrate the use case a scenario is presented in ???, where a drone is loitering around an accident site. The operator can provide a ROI to the interface as a GPS point, but will have to go through the process of choosing and uploading a new point each time.

1.3 Research Motivation

With the expansion of newer, low-latency network technologies like 5G, the use case for remote control has grown stronger. With everything from connected cars, remote surgery or log lifting, the uses of real-time remote-control are many.

In the research field of Quality of Experience the user experience of remote control is studied, and common objects of study are the effects of different network parameters like latency, jitter or image quality. Previous QoE research have been on applications like log lifting [17] and excavation equipment [15], and the highly dynamic environment of a drone flying above the sea can make an interesting study object.

With the promise of real-time remote control over 5G, the systems of tomorrow will be heavily reliant on the uptime and latency of the network, making these systems exposed to vulnerabilities should the network experience high contingency or failure. Therefore, Having a system that can operate in a degraded state, i.e. with high latency, is a necessity. Furthermore, when 5G is more widely adopted and stable, Internet Service Providers (ISPs) might offer plans with different latencies, giving the user a choice of how much they want to pay for latency that their particular application needs.

1.4 Scope

The scope of the thesis work is to develop manual gimbal control using the ArduPilot firmware, running on a flight controller connected to a servo gimbal. This prototype is then to be used as a test bed for a QoE experiment evaluating the effects of latency on the operator's experience and task performance. If time allows the interface will be integrated into the existing software stack of the EOS project, serving as a proof of concept for manual gimbal control during flight.

Chapter 2

Background

In this chapter a brief background will be given of the Swedish Sea Rescue Society along with the drone project that this work is a part of. Then, the research field of Quality of Experience will be introduced, along with the related previous work.

2.1 Swedish Sea Rescue Society

As can be read on their website [11], the Swedish Sea Rescue Society (SSRS) is a non-profit organization that was founded at a conference in Stockholm in 1906 due to Sweden receiving criticism of its' poor sea rescue. Today, it is a foundation with 40 employees and over 143 000 members, and with 2400 volunteers manning their 260 rescue vessels, they carry out around 90% of all sea rescues in Sweden all year around.

2.1.1 Innovation

As stated in their statutes [12], the mission of SSRS is not solely to carry out these rescue missions but to also innovate and collaborate in the area of maritime rescue as well as other aiding activities in society as a whole. As a result of this, the drone project that this thesis is a part of has been conceived along with other innovation projects such as foiling rescue boats and improved rescue vehicles [10].

2.1.2 Drone Project: Eyes On Scene

This thesis work is part of a project called Eyes-on-Scene, where a task-specific drone and mission control software has been developed to meet the requirements of sea rescue missions. In this section the drone and surrounding system will be briefly described.

The drone uses a fixed-wing frame and is designed for two modes: sprinting and loitering, which allows for short response time while also being able to stay in the air for a long time. The only equipment it carries is a single camera gimbal with adjustable roll, pitch and yaw angles. The current system uses a launchpad for takeoff and the idea is that the plane will be able to land on a rescue vessel or in the water, allowing it to be picked up.

A web interface for mission control has also been developed in the project. Currently, the interface allows flying to and loitering around waypoints given on a map which also displays both other naval and aerial traffic in real time. The map also shows the regulatory zones in which one is not allowed to fly, i.e. close to airports or military zones. The interface provides video from the drone camera which is recorded and accessible after the mission.

2.2 Quality of Experience

Quality of Experience is an emerging research field that is concerned with the user's experience in multimedia systems. It is an inherently multidisciplinary field that has close ties to the field of User Experience (UX) and Human-Computer Interaction (HCI). QoE is also closely related to the field of Quality of Service (QoS), which is concerned with the technical aspects of a system, but aims instead at evaluating the subjective experience through user experiments rather than qualitative measurements.

In the white paper [14], Brunström et al. make a working definition of Quality of Experience:

Quality of Experience (QoE) is the degree of delight or annoyance of the user of an application or service. It results from the fulfillment of his or her expectations with respect to the utility and / or enjoyment of the application or service in the light of the user's personality and current state.

2.3 Previous Work

W. Tärneberg et al. [17] present a QoE study with industrial equipment on an excavation site where experienced operators controlled their usual equipment remotely at different latencies.

K. Brunnström et al. [15] performed a study on log lifting using a head-mounted display system. In the study, latency is introduced both in the display's response to movement as well as the controls. The display delay was found to have a strong effect on nausea, but an observable effect on controller latency could only be found at latencies above 800 ms.

As part of the Eyes-On-Scene-project, a study was performed at Chalmers by Grote et al. [16] where the same emergency call was performed with and without images from a UAV at the emergency site. The study showed that imagery from the accident gave the rescue personnel a sense of control before arriving at the scene when knowing what they were going to face. The crew was also faster at locating a person or object when having aerial footage of the scene.

Chapter 3

Test Beds

This chapter will go through the technical details of the test beds and their components as were used in the experiments. First, all the hardware will be presented. Then, the pre-existing software along with the software implemented specifically for the thesis work will be detailed along with some discussion about design choices made.

3.1 Hardware

In this section the hardware used in the experiments is presented.

3.1.1 Fixed-Wing Drone

As part of the EOS-project a fixed-wing drone has been developed. It belongs to the SSRS and is designed together with plane manufacturer Smartplanes. The plane can be seen in 3.1.

The components relevant for this thesis work are:

Flight Controller: Pixracer R15 The flight controller is the brain of the aircraft and houses components and connectors relevant to in-flight operations such as servos for control surfaces or other equipment, GPS and accelerometer.

Companion Computer: Raspberry Pi 4 The Raspberry Pi serves as the onboard computer, commonly known as the "companion computer" in the context of drone hardware. It is connected to the flight controller via USB and makes it possible to relay messages sent over ethernet instead of radio, allowing to connect to a ground station over the Internet.

4G Modem Sim-card modem allowing access to the mobile network.



Figure 3.1: The SSRS fixed-wing drone. The left picture is the inside of the drone housing with the internals visible and the right picture shows the housing mounted on the airframe.

Gimbal The gimbal is designed and manufactured by Fredrik Falkman at the SSRS. It is a 3D-printed design that has three degrees of freedom made possible by three servos connected to drive belts that control each axis of the camera. The servos are connected to the flight controller which also provides them with power. When mounted the servo-housing is inside the aircraft while the part with the mounting plate for the camera module sticks out on the bottom of the drone. In 3.2 the schematics of the gimbal can be seen.

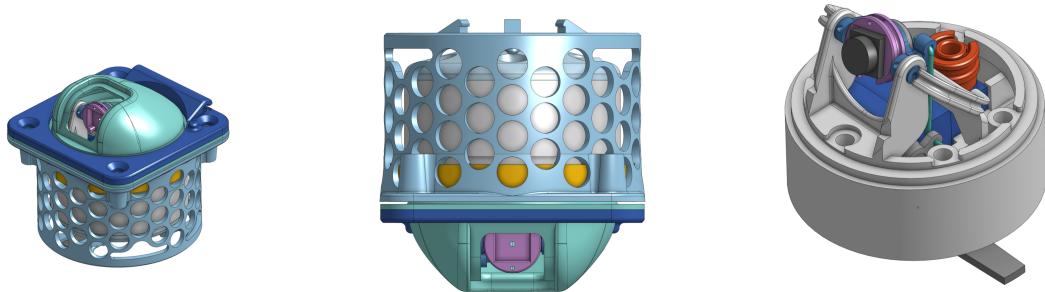


Figure 3.2: Schematics of the drone gimbal. The middle and left picture shows the gimbal with its full housing. To the right the housing has been removed and a mockup camera module has been inserted in the mounting plate.

Raspberry Pi Camera Module v2 The camera module is a small camera that is mounted inside on a flat surface on the gimbal. The camera connects to the module board which

is then connected to the Raspberry Pi with a ribbon cable. It is capable of recording 1440x1080 video at 30 fps and has a 4:3 aspect ration.

3.1.2 Gimbal Package

As the entire drone was not needed to conduct the QoE experiment, a Raspberry Pi, flight controller and gimbal were mounted inside a 2L Ikea SmartStore. The boards were mounted inside the box with velcro tape and the gimbal was mounted in a cutout in the bottom of the box. Pictures of the box can be seen in figure 3.3.



Figure 3.3: The box that houses the drone hardware for the experiment. The left picture shows the box from above, from the left the following components are visible: flight controller, gimbal with camera module, power supply and Raspberry Pi. The right image shows the box from below.

3.1.3 Hoverboard robot

In order to introduce movement in QoE experiment, a modified hoverboard was used. The hoverboard has four wheels and could be controlled via joystick inputs or coordinates. Its manouevring was assisted by real-time lidar-mapping of the room. The hoverboard can be seen in 3.4

3.1.4 Apriltag

An apriltag is a type of low-resolution QR-code often used in robotics to estimate orientation and position of robots or objects with the help of computer vision. It was fastened on top of the hoverboard as high as possible without blocking the lidar. The apriltag can be seen mounted on the hoverboard in figure 3.4.

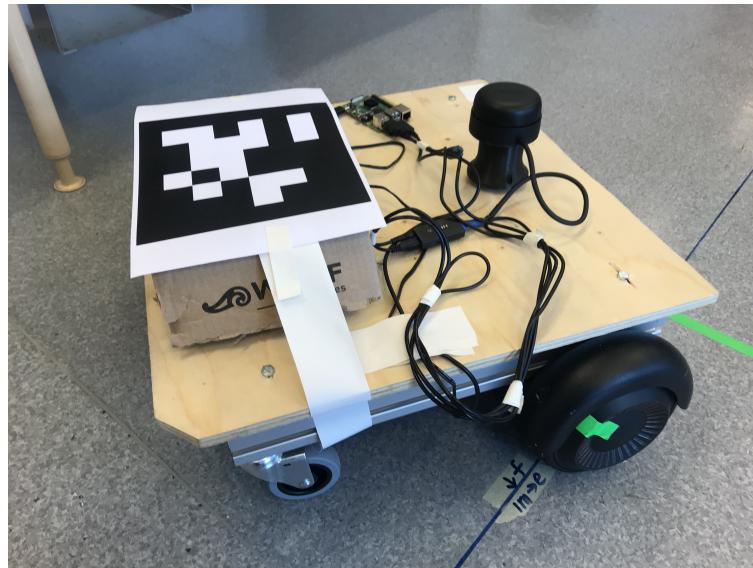


Figure 3.4

3.2 Software

In this section the software used in the experiments will be detailed.

3.2.1 Pre-existing Software

Firmware: ArduPlane The firmware running on the flight controller is called ArduPlane, which is part of the open-source autopilot software suite ArduPilot that enables the creation of unmanned, autonomous vehicles [2]. It is developed by a community of developers and is available on GitHub [1].

Protocol: MAVLink As can be read on their website [5], MAVLink is a lightweight protocol suited for communication with drones and between drone components. It has a byte overhead of 14 bytes and allows for concurrent communication between up to 255 systems.

Software: MAVProxy MAVProxy [6] is a software running on the companion computer whose task is to relay MAVLink messages from the control station to the flight controller. In the test bed it relays messages between the ethernet port (application) and the serial port (drone box).

UV4L UV4L is a video streaming server that supports real-time communication protocols and video encodings [13]. Its function in the test bed is to stream the local video feed to the web.

Janus Server As described on their webpage [4], Janus is a plugin-based software that helps establish WebRTC connections. In this thesis it is used to establish the connection between the UV4L server and the web browser, enabling the peer-to-peer connection between the two devices.

Hoverboard software The hoverboard has software which allows it to be controlled by a game controller or follow a set of points using coordinates for which it uses its Lidar and a SLAM-algorithm. This makes it possible to have the hoverboard run in a programmed route with little error.

3.2.2 Gimbal Control Interface

A program displaying the video feed from the camera while taking joystick inputs from the operator was implemented in Python. The software took the joystick inputs from a PlayStation-controller and could send messages to the flight controller updating the position of the gimbal at a desired frequency and delay. The program could also record the video feed. A visual overview of the system is shown in Figure 3.5.

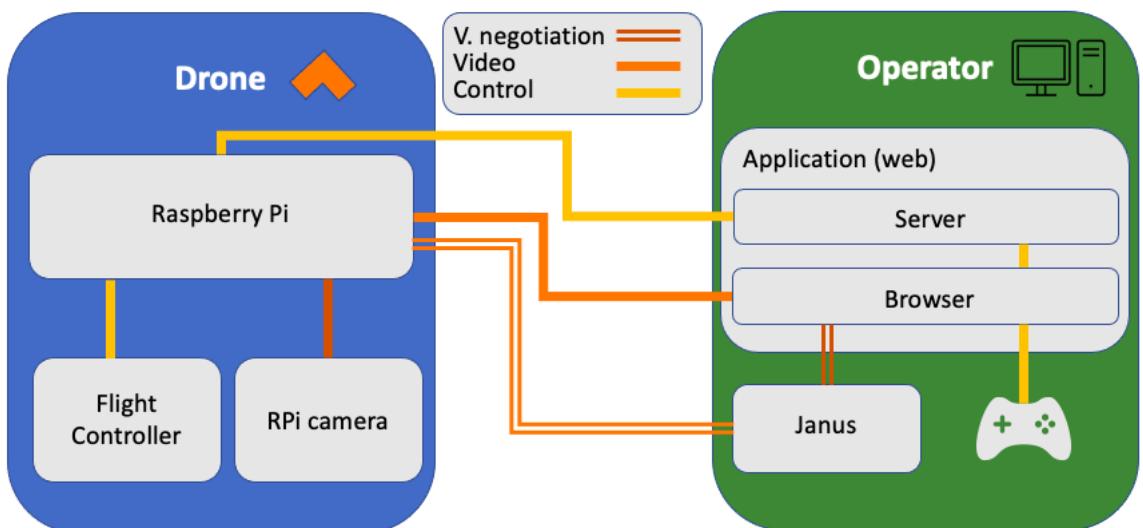


Figure 3.5: An overview of the entire system. On the left, the drone and its components including the RPi, RPi camera module and the flight controller are shown. On the right, the ground station and its' components including the web application, Janus server and the controller are shown. The colour of the connections represent what data they transfer, where orange is video, orange double-line is video negotiation and yellow is control.

The application was run with Python version 3.9.16 and used the following libraries:

pymavlink (v. 2.4.37) An interface to communicate with MAVLink devices, available at [9]. With the library a port can be set to send and receive MAVLink messages.

Evdev v. (1.6.1) A python library for interrupt-driven input devices. It is used for reading input from the Playstation-controller. Available at [3].

OpenCV (v. 4.7) Open Computer Vision, a very popular library for computer vision tasks. In the test interface it is used for displaying and recording the video feed. Available at [8].

Multiprocessing This standard module was used in order to run the video tasks in different processes to not interfere with the control. The video frames were accessed and displayed by one process and then feeded to another process saving each frame. For more information on the library see [7].

Delay was introduced in the controls by entering commands in a queue along with a timestamp, with the program waiting until timestamp was reached before sending the command. The following algorithm was used to schedule the sending of commands according to the delay:

```
while running do
    timestamp, pitch, yaw = queue.get() // blocking call
    t0 = time.now
    diff = timestamp + delay - t0
    if diff > 0 then
        time.sleep(diff)
    end if
    send_command(pitch, yaw)
end while
```

The inherent delay of the system was measured by recording a stopwatch, the controller and the video feed of the drone simultaneously. The screen and the controller were then filmed on a phone in 120 FPS to measure the time between a control input to the video feed updating. To measure the network delay a simple ping-measure was taken.

The network latency was found to be below 1 ms. The delay of the video feed alone was between 100-150ms while the full controller to video feedback delay was around 400ms.

Although the sources of the delay were not examined in any further depth there are a few probable causes of the delay:

Message frequency While the internal model of the gimbal's orientation was updated by interrupts, the message rate to the flight controller was set to a maximum of 20Hz.

Video encoding Some video encodings are faster at the cost of quality and vice versa. Although not the fastest, MJPEG was chosen as it was simple to integrate with OpenCV and had acceptable latency and video quality.

MAVProxy Each message must be taken from the ethernet interface of the Raspberry Pi and be sent on the serial interface to the flight controller. It must then be processed before the servos can be updated.

3.2.3 Result Script

In order to get results from the recorded video of each subject a script was implemented that was run on each trial, measuring the error of the operator in each video frame and saving it to a file.

3.2.4 Auto-follow mode

As a proof of concept, an auto-follow mode was developed which allowed the gimbal to follow the hoverboard. This was made using the object detection algorithm from the result script,

but in real-time coupled with a PI-controller which had the center of the image as reference value. Due to time limitations this was not investigated any further.

3. TEST BEDS

Chapter 4

Experiments

This chapter will describe the setup, procedure and evaluation of the two experiments that were performed.

4.1 QoE Lab Experiment

This section describes the procedure of the experiment aimed at evaluating the effect of latency on the gimbal operator using the hoverboard robot.

4.1.1 Task

The task was to keep the center of the screen, marked with a red dot (as seen in 4.4), aligned with the center of the apriltag laying on top of the hoverboard (shown in 3.4) while moving in an ellipse.

4.1.2 Experimental Setup

The experiment was set up in a lab where the subject and the conductor sits at a desk with a high shelf behind them obscuring the view towards the rest of the lab. Behind the shelf the robots used in the experiment are located. An illustration of the lab setup is provided in figure 4.1.

An image of the hoverboard and the camera is shown in image 4.2.

Subject selection and recruitment

The subjects were recruited mostly from the conductor's network of contacts, as in order for the subjects to be covered by insurance during the experiment, the subjects had to be either

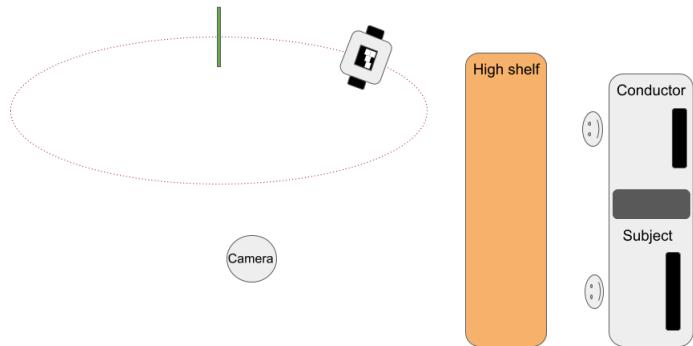


Figure 4.1: An overview of the experimental setup. To the right the experiment conductor and subject by a desk facing right. On the left side, behind the shelf, the test bed including the camera and the hoverboard are located. The camera is mounted at a height of 2.1 meters with the gimbal facing down.

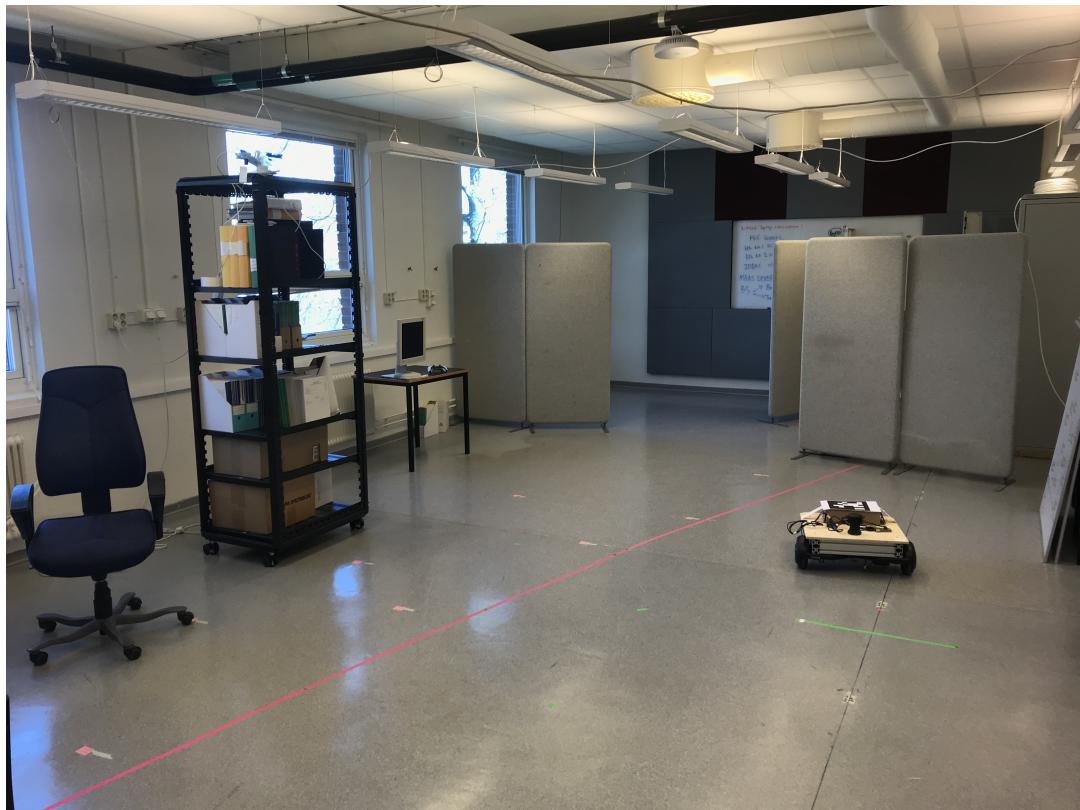


Figure 4.2: A picture of the test bed. On the left side there is a server rack with the drone box mounted on top. The hoverboard can be seen on the right side of the picture.

be students or employed by Lund University. Some of the participants came spontaneously while others booked a time in a spreadsheet provided by the conductor.

There were 10 test subjects, 4 women and 6 men, with an average age of 26.1 years where the youngest was 23 and the oldest 37. Furthermore, they were questioned about their eyesight, handedness as well as experience with FPV-games and flying drones.

There was no compensation for the subjects other than home baked goods and a hot beverage, which were offered at the beginning of each experiment.

Lab Conditions

The lab is located in the E-building at LTH, Lund University. It can also be mentioned that the windows in the room face north, so there was no risk of direct sunlight during the experiments. The computer screen was 24' and the application was run on a machine with the following specifications:

CPU Intel Core i7-4770 @ 3.40GHz, 4 cores 8 threads

RAM 16 GB

GPU Mesa Intel HD Graphics 4600 (HSW GT2)

4.1.3 Experimental Procedure

In this section each of the steps will be described in more detail. The experimental procedure is summarized in the table 4.1.

Table 4.1: Task Timeline

Time (approx.)	Task	Instruction
3	Coffee and cookies	
5	Introduction	Tell the subject about what the study is about. Gather personal data + consent
2	SSQ	
3	Experiment walkthrough	What is going to happen Description of task Where to aim exactly on the apriltag
2	Warm-up	The subject gets to try the setup for one lap
10	Tests	5 tests with added latencies [0, 200, 300, 400, 500] in random order
2	SSQ	
3	Interview	
30	Total	

Introduction

The experiment starts with the subject being seated at the desk to read through the consent form. After this was signed the conductor clarified that the subject was free to ask any question during the experiment and could choose to interrupt it without declaring any reason.

Then, the subject fill out the background information form where they also get to choose a four digit code which from that point is the only identifier of the subject and it's results. If the subject booked a time slot in the spreadsheet, they were at this point removed from it.

Experiment Walkthrough

The task was explained to the subject while a printed apriltag with a cross on it was shown, marking the spot where to aim on the tag. The particular apriltag printed for this experiment had two white and two black squares meeting in the middle, making the center easier to identify. The apriltag shown to the subject is shown in figure 4.3.

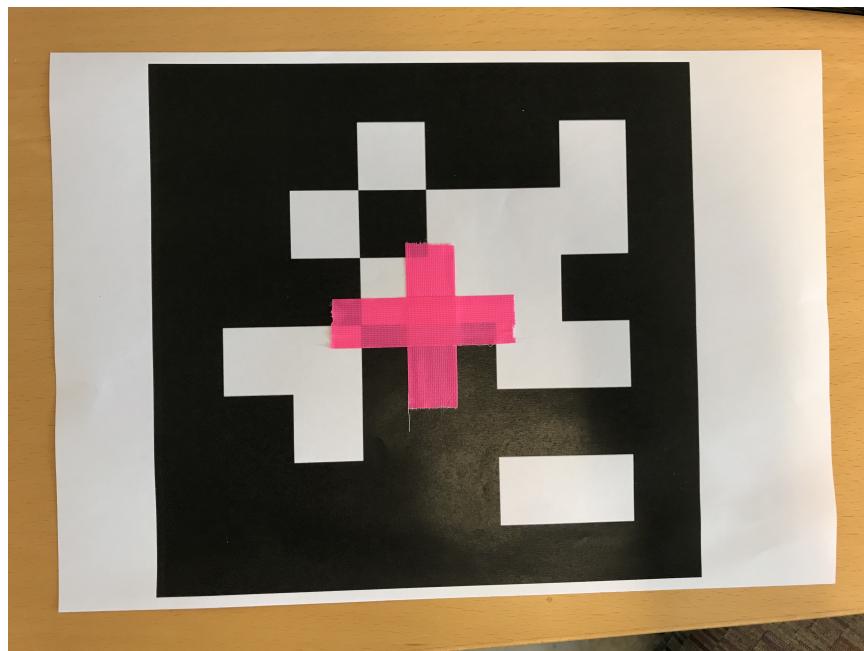


Figure 4.3: The apriltag used to show the subject where to aim.

The subject was then given a chance to try the system without any delay, and was given the controller and presented with the view, shown in figure 4.4. The hoverboard was started and ran for one lap in the same path that it would run during the test and the subject could try and follow it.

Test Procedure

After the task description and warm-up the tests commenced. The added latencies tested were 0, 200, 300, 400 and 600 and the order in which they were given to the subject was random. At the beginning of each trial the interface was restarted with one of the latencies induced in the system. The subject was then asked to put the red marker in the middle of the screen in the middle of the apriltag. When the subject confirmed that they were ready to start the conductor started the hoverboard. The subject followed the hoverboard for two laps, which took about 2 minutes in total.

After the laps were completed the interface was shut down and the subject was asked to rate the system by answering four different questions on a scale from 1 to 5. When the subject had answered the questions the next trial was started.

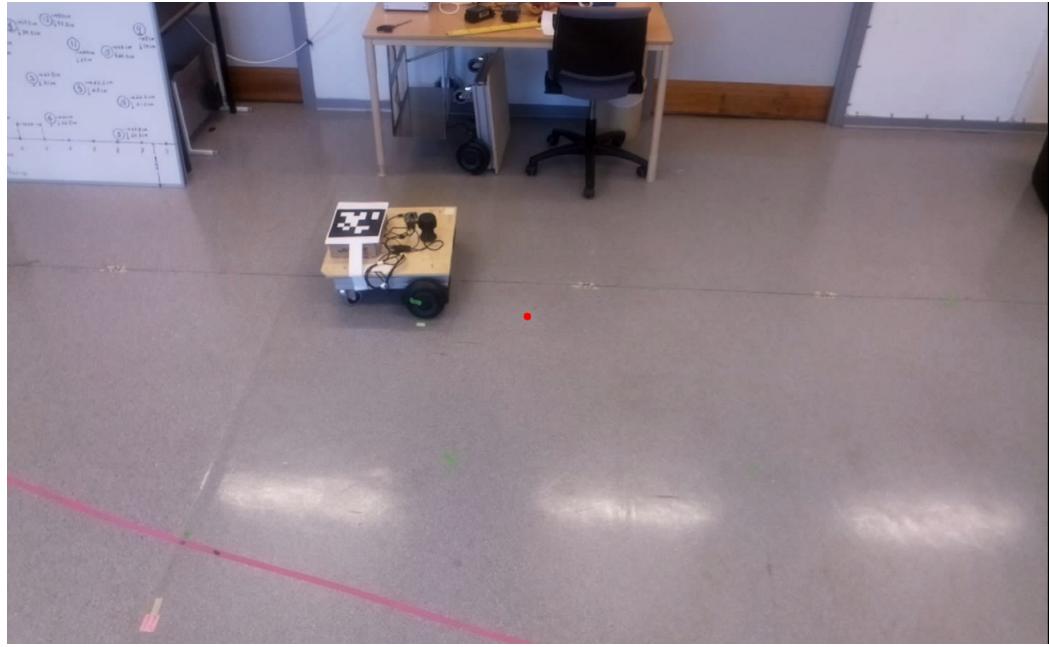


Figure 4.4: The interface that the subject was presented with. The red dot in the middle of the screen is overlayed on top of the displayed video so that the subject has a better reference of the center of the screen.

At the end of the fifth and last trial the subject got to fill in the sickness questionnaire was also asked a few interview questions by the conductor.

4.1.4 Evaluation

The following two sections will describe the quantitative and qualitative methods that were used to evaluate the system.

Quantitative Evaluation

To measure the performance of each subject, a script using a detection algorithm was used on the recorded video which for each frame gave the distance between the red dot and the center of the apriltag. Pythagoras' theorem was used for calculating the pixel distance between the two points and is shown in equation 4.1. A frame from the detection script can be seen in figure 4.5.

$$d = \sqrt{(x_{hoverboard} - x_{center})^2 + (y_{hoverboard} - y_{center})^2} \quad (4.1)$$

Qualitative Evaluation

After each trial the subject answered a form with four questions about the system on a five degree scale, labeled 'Terrible' to 'Excellent'. The questions are shown in table 4.2.

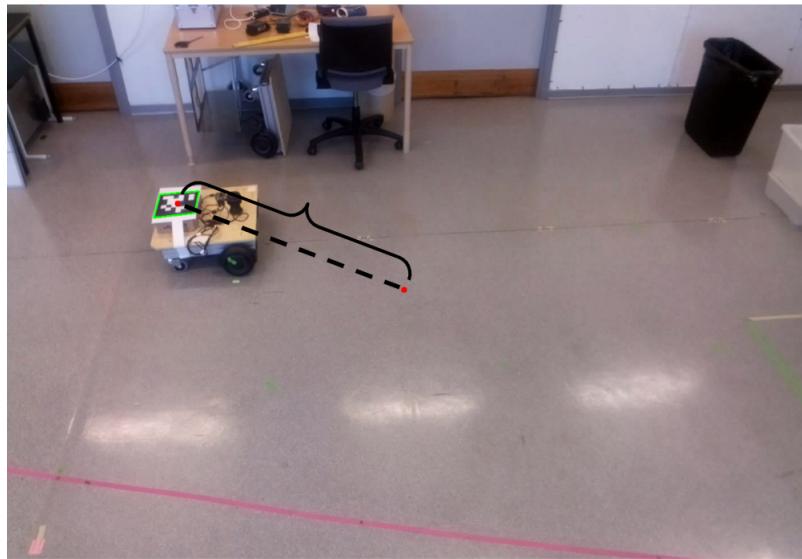


Figure 4.5: Example of a result from the detection algorithm. The green square outlines the result of the detected object and the red dot in the middle of its center. The error is the distance between the two dots as outlined by the black indicators.

How would you describe the controllability of the system?
Terrible Poor Fair Good Excellent
How would you rate your ability to perform the task?
Terrible Poor Fair Good Excellent
How would you rate your experience of the system?
Terrible Poor Fair Good Excellent
How would you rate your impression of the system as a whole?
Terrible Poor Fair Good Excellent

Table 4.2: The questions asked for each delay after each trial.

Right after the last trial the subject filled out the SSQ once more. Then, a brief interview was conducted with questions about the system and experience in general. The questions can

be seen in table 4.3.

1. What was your general experience of controlling the camera?
2. Did you experience any difference in the controls between the trials?
3. Was there any part of the track that was harder than any other?
4. Do you think the system would have been usable with the worst experimental conditions?

Table 4.3: The interview questions asked by the conductor after the last trial.

4.2 Known Sources of Error

4.2.1 Ecological validity

The future operators of the system are to be the volunteers of SSRS, which have a wide variety in age and background. Thus, there are three main concerns with the ecological validity of the study. Firstly, the subjects were all either students or researchers at Lund University. Secondly, the average age of the subjects was low. Lastly, the sample size was only ten people.

The last concern is probably least important when it comes to the validity of the study, as the low number of participants could be compensated for by having each subject provide more data. Each subject generated around 13 500 data points in five trials with two laps each, resulting in about 135 000 data points in total. This is likely to have had a positive effect on the validity of the results.

The second and third concerns are more difficult to compensate for, as only employees or students at the university could be used as subjects due to the university insurance. There are of course people of different ages and backgrounds working and studying at the university, but due to time restraints as well as lack of compensation for the subjects, most came from the conductors network of contacts.

4.2.2 Frame detection accuracy

The detection algorithm used to calculate the distance from the center of the apriltag to the center of the screen could sometimes during fast camera movements not detect the apriltag. This resulted in a NaN value at that particular frame. The average percentage of frames detected where 85%, with a minimum of 80% and a maximum of 89%. The percentage of frames detected was not found to be correlated with the delay and are plotted for each subject and delay in figure 4.6.

To make for better averages over different trials, all NaN values, i.e frames where no apriltag was detected in the frame, were linearly interpolated. To illustrate this one result is shown before and after interpolation in 4.7.

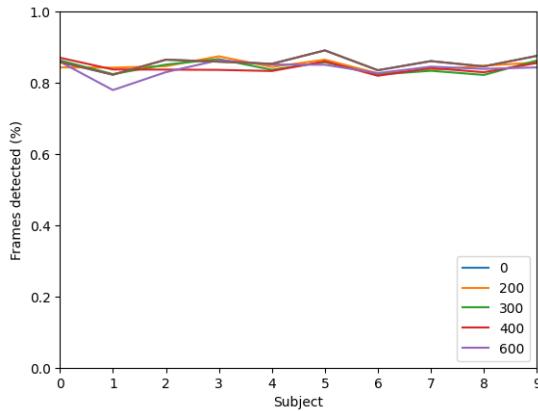


Figure 4.6: Form answers averaged for all subjects for each delay.

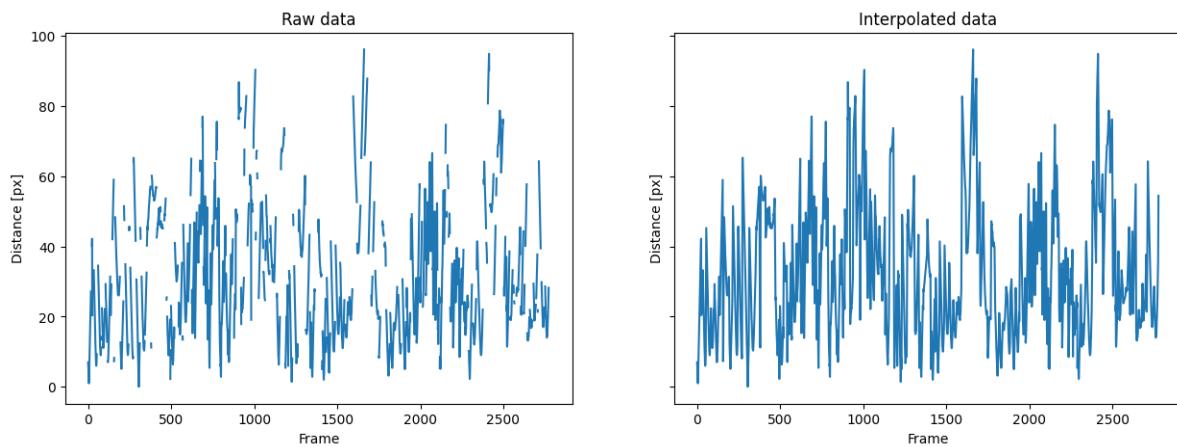


Figure 4.7: Results for one trial: to the left the result with NaN-values and to the right interpolated result.

4.2.3 Difficulty of the task

The center of the apriltag did not have a very clear center which could have made it difficult for the subject to complete the task. To remedy this, a printed apriltag where the center was marked out, but it could be that the subject had difficulty identifying the exact center during the trial.

4.3 Field testing

This experiment aimed at running the gimbal interface on the real drone in-flight. The SSRS have a permit to operate drones beyond visual line of sight within a delimited area of the archipelago of Gothenburg. Because of this, the drone can be flown longer distances while all communication going over the mobile network.

4.3.1 Goals

The main goals of the field testing were:

- confirm that the gimbal control did not interfere with the flight control software
- evaluate suitability of the controls in a dynamic environment
- compare the manual controls to the already existing ROI-mode

4.3.2 Preparations

For the experiment the gimbal control interface used earlier was modified in order to run with the software and video streaming configured on the actual drone. The video and recording component was entirely removed and the IP address which the program connected to was changed to that of the 4G modem. In order to compare the new controls to the already existing ROI-mode, a button on the controller was mapped to switch between joystick-mode and ROI-mode, using pre-programmed GPS-points for the drone to look at.

4.3.3 Procedure

A flight of roughly 30 minutes with 17 waypoints and 2 loiter (circling) points was performed. For the waypoints the camera was held in manual mode for the most of the time, enabling the camera operator to look around using the joystick. The latency of the video was estimated to be between one and two seconds, which gave some delay to the controls. However, the gimbal was still controllable and the camera operator could still look around and point the camera at objects of interest.

For convenience, a button on the control was mapped to switch the gimbal control between manual mode and ROI-mode available in the firmware. This made it possible to easily switch between the two.

During loitering the camera was difficult to control manually as the plane was moving around a lot. When going around in a circle the plane also moves a lot in the yaw and roll axes, making it very difficult to compensate manually with the joystick. The ROI-mode was better suited for this, although there were some inconsistencies in the altitude data for the points which caused the camera to point at a point below the surface of the earth. The point was adjusted by raising the altitude provided to the GPS-point.

Chapter 5

Results

This chapter will go through the main findings of the experiment. The quantitative measurements will first be presented followed by the results from the forms and interviews.

5.1 Quantative measurements

In 5.1 the mean and standard deviation of the different delays are presented. The values are also visualized in plot 5.1. A clear trend can be seen where both the mean and standard deviation increase with the delay. An interesting feature of the data is that the mean after 400ms starts to increase at a non-linear rate, while the standard deviation has a more or less linear increase all throughout the added delays.

Delay	Mean	Std. deviation
0	29.804138	18.030196
200	35.234192	22.347311
300	41.538190	26.805780
400	46.048449	28.718249
600	61.811312	37.226394

Table 5.1: Table caption goes here.

In figure 5.2 the empirical cumulative distribution functions (ECDF) for the results of each delay are shown. The x-axis represents the value which the entry has while the y-axis represents the percentage of entries that are less than or equal to the x-axis value. The increase in both average and mean can be clearly observed as the delay increases. It can also be deduced that the distance between the center of the ECDF distributions for 0ms and 200ms is much

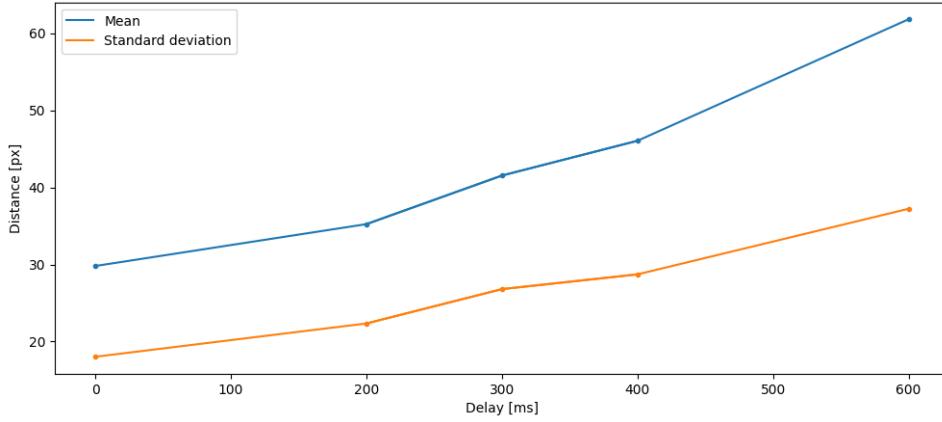


Figure 5.1: Pixel error average and standard deviation for all delays. Delays 100 and 500 are added on the x-axis for correct scale.

smaller than that between the curves for 400ms and 600ms. This could indicate that the added 200ms added delay had a larger impact at 400ms than at 0 ms.

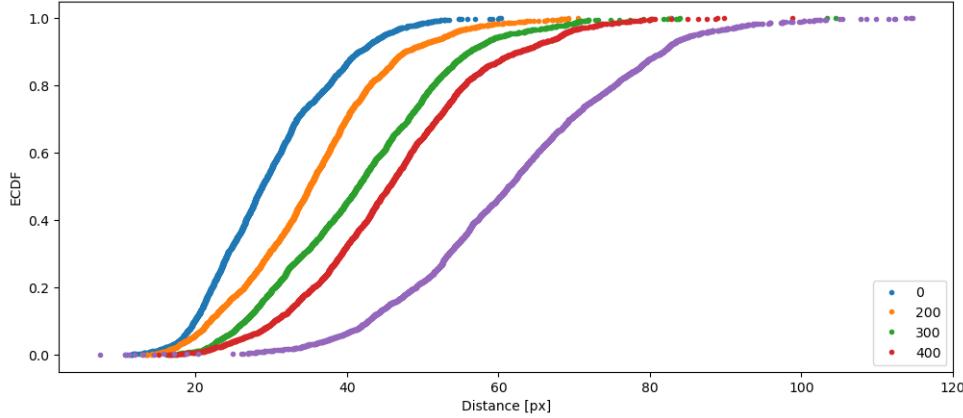


Figure 5.2: ECDF of the total distance for all subjects. The curves show the distribution of the error for each delay. A larger inclination suggests a larger spread and the more to the right the curve is, the larger the average error.

In figure 5.3 all the values from the trials have been averaged and the position of the hoverboard has been superimposed. In this plot one can see that the final corner of the track is the most difficult, and that the very end and early beginning of the track are the easiest.

To see the performance over the course of the hoverboard's track for each delay, the trials of all subjects were summarized on a particular delay alongside the position of the hoverboard. This is shown in 5.4, where all trials with 0ms and 600ms added delay are averaged on each frame. The more transparent blue and green lines represent the average over all trials while the darker lines are a rolling average of the data. The red line represent the position of the hoverboard. It can be seen that there is a

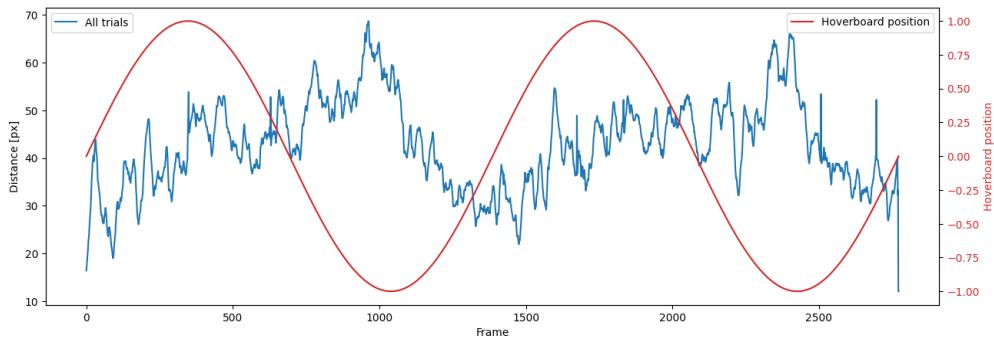


Figure 5.3: All trials averaged with hoverboard position superimposed. The error steadily increases up until the end of the last corner where it decreases rapidly.

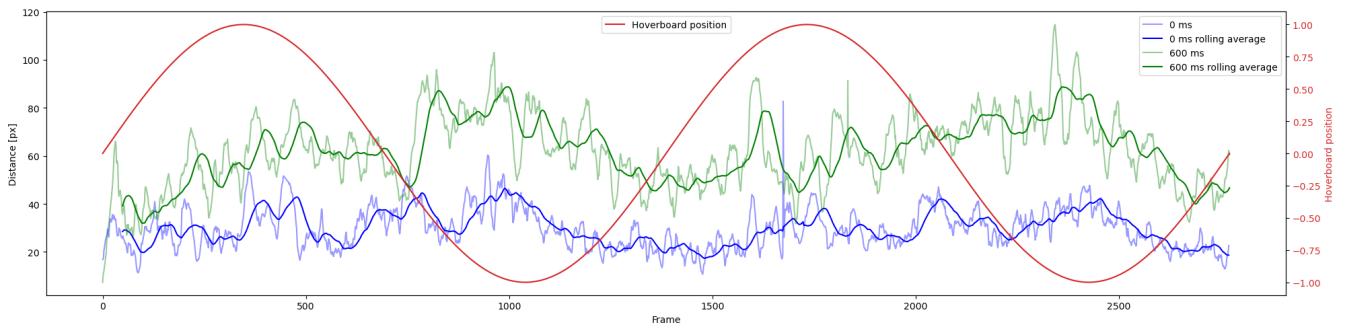


Figure 5.4: Total trials for latencies 0 and 600 averaged over all subjects.

5.2 Forms and interviews

In 5.5 the answers from the questionnaires filled out after each trial are averaged over all subjects and visualized as a grouped bar chart. It is important to note that the jumps in delay are not uniform. A monotonous decrease can be seen on all questions where higher delays give lower scores on average. On the second question, namely "ability to control task", the results for 200ms, 300ms, and 400ms are similar. The jumps from 0ms to 200ms and 400ms to 600ms seem to be much more impactful on the subject's perceived ability to perform the task than the jump from 200ms to 400ms.

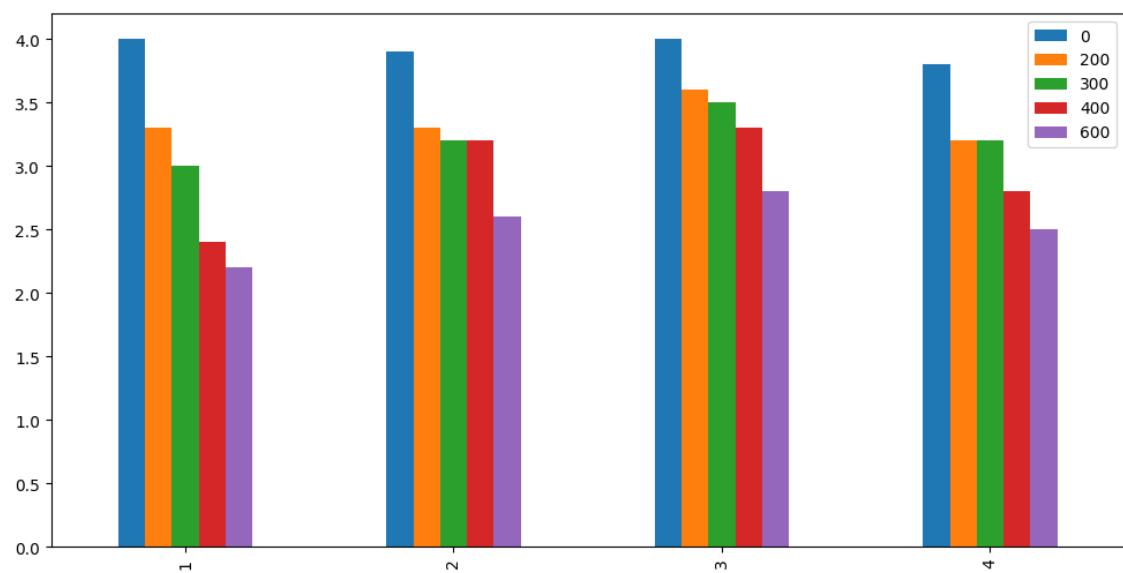


Figure 5.5: Form answers averaged for all subjects for each delay.

Chapter 6

Discussion

As the goal of this thesis work was to develop an application that could be used in a controlled experiment as well as in actual flight, there were some design choices that were taken during development. This section will briefly discuss the reasoning behind the design choices taken.

6.1 Video streaming

In the real drone system the video feed is done with WebRTC. This is a modern peer-to-peer video protocol well suited for the application of video from IoT, and the intention was to use this in the test bed as it had a low latency and also was accessible through the browser. However, it proved more difficult to access as other video processing software had to be used to access the raw video feed, which was needed for adjusting image parameters as well as for detection and recording of the images. In the end the MJPEG format was chosen, as it had acceptable delay and could be integrated only by providing a web address to the OpenCV video capture module.

6.2 Native vs. web

The application that the SSRS currently use for mission control is hosted on the web, and as such the test bed was initially developed as web-application. This proved difficult, as the developer was not experienced with multithreaded web applications necessary for the tasks that were to be performed simultaneously, i.e. control, recording, detection. As the web approach proved too complex, the test bed was implemented as a native python application.

6.3 Sources of error

6.3.1 Frame detection accuracy and interpolation

6.4 Experiment execution

Chapter 7

Conclusion

Chapter 8

Future Work

References

- [1] Ardupilot github, ardupilot. <https://github.com/ArduPilot/ardupilot>. Accessed: 2023-02-13.
 - [2] Ardupilot org, ardupilot. <https://ardupilot.org/>. Accessed: 2023-02-13.
 - [3] evdev, pip. <https://pypi.org/project/evdev/>. Accessed: 2023-05-02.
 - [4] Janus, meetecho. <https://janus.conf.meetecho.com/>. Accessed: 2023-03-10.
 - [5] Mavlink. <https://mavlink.io/en/>. Accessed: 2023-02-13.
 - [6] Mavproxy. <https://ardupilot.org/mavproxy/>. Accessed: 2023-03-28.
 - [7] Multiprocessing, python. <https://docs.python.org/3/library/multiprocessing.html#module-multiprocessing>. Accessed: 2023-05-02.
 - [8] opencv, pip. <https://pypi.org/project/opencv-python/>. Accessed: 2023-05-02.
 - [9] pymavlink, github. <https://github.com/ArduPilot/pymavlink>. Accessed: 2023-02-13.
 - [10] Surtsey innovation projects. <http://www.surtsey.org/>. Accessed: 2023-02-03.
 - [11] Swedish sea rescue society, information in english. <https://www.sjoraddning.se/information-english>. Accessed: 2023-02-01.
 - [12] Swedish sea rescue society, statues. https://www.sjoraddning.se/sites/default/files/stadgar_ssrs_antagna_2018.pdf. Accessed: 2023-02-03.
 - [13] Uv4l, linux projects. <https://www.linux-projects.org/uv4l/>. Accessed: 2023-03-10.
-

- [14] Kjell Brunnström, Sergio Ariel Beker, Katrien de Moor, Ann Dooms, Sebastian Egger, Marie-Neige Garcia, Tobias Hossfeld, Satu Jumisko-Pyykkö, Christian Keimel, Mohamed-Chaker Larabi, Bob Lawlor, Patrick Le Callet, Sebastian Möller, Fernando Pereira, Manuela Pereira, Andrew Perkis, Jesenka Pibernik, Antonio Pinheiro, Alexander Raake, Peter Reichl, Ulrich Reiter, Raimund Schatz, Peter Schelkens, Lea Skorin-Kapov, Dominik Strohmeier, Christian Timmerer, Martin Varela, Ina Wechsung, Jun-yong You, and Andrej Zgank. Qualinet White Paper on Definitions of Quality of Experience, March 2013. Qualinet White Paper on Definitions of Quality of Experience Output from the fifth Qualinet meeting, Novi Sad, March 12, 2013.
- [15] Kjell Brunnström, Elijs Dima, Tahir Qureshi, Mathias Johanson, Mattias Andersson, and Märten Sjöström. Latency impact on quality of experience in a virtual reality simulator for remote control of machines. *Signal Processing: Image Communication*, 89:116005, 2020.
- [16] Benjamin Grote and Leo Granholm. Evaluating the effect of livestreamed video from accident sites in maritime sar using uav. 2022.
- [17] William Tärneberg, Omar Hamsis, John Hedlund, Kjell Brunnström, Emma Fitzgerald, Andreas Johnsson, Viktor Berggren, Maria Kihl, Akhila Rao, Rebecca Steinert, and Caner Kilinc. Towards intelligent industry 4.0 5g networks: A first throughput and qoe measurement campaign. 2020.

Appendices

Appendix A

About This Document

The following environments and tools were used to create this document:

- operating system: Mac OS X 10.14
- tex distribution: MacTeX-2014, <http://www.tug.org/mactex/>
- tex editor: Texmaker 5.0.2 for Mac, <http://www.xm1math.net/texmaker/> for its XeLaTeX flow (recommended) or pdfLaTeX flow
- bibtex editor: BibDesk 1.6.3 for Mac, <http://bibdesk.sourceforge.net/>
- fonts `cslthse-msc.cls` document class:
 - for XeLaTeX: TeX Gyre Termes, TeX Gyre Heros, TeX Gyre Cursor (installed from the TeXLive 2013)
 - for pdfLaTeX: TeX Gyre font packages: tgtermes.sty, tgheros.sty, tgcursor.sty, gtxmath.sty (available through TeXLive 2013)
- picture editor: OmniGraffle Professional 5.4.2

A list of the essential L^AT_EX packages needed to compile this document follows (all except `hyperref` are included in the document class):

- `fontspec`, to access local fonts, needs the XeLaTeX flow
- `geometry`, for page layout
- `titling`, for formatting the title page
- `fancyhdr`, for custom headers and footers
- `abstract`, for customizing the abstract

- `titlesec`, for custom chapters, sections, etc.
- `caption`, for custom tables and figure captions
- `hyperref`, for producing PDF with hyperlinks
- `appendix`, for appendices
- `printlen`, for printing text sizes
- `textcomp`, for text companion fonts (e.g. bullet)
- `pdfpages`, to include the popular science summary page at the end

Other useful packages:

- `listings`, for producing code listings with syntax colouring and line numbers

A.1 Page Size and Margins

Use A4 paper, with the text margins given in Table A.1.

Table A.1: Text margins for A4.

margin	space
top	3.0cm
bottom	3.0cm
left (inside)	2.5cm
right (outside)	2.5cm
binding offset	1.0cm

A.2 Typeface and Font Sizes

The fonts to use for the reports are **TeX Gyre Termes** (a **Times New Roman** clone) for serif fonts, **TeX Gyre Heros** (a **Helvetica** clone) for sans-serif fonts, and finally **TeX Gyre Cursor** (a **Courier** clone) as mono-space font. All these fonts are included with the TeXLive 2013 installation. Table A.2 lists the most important text elements and the associated fonts.

A.2.1 Headers and Footers

Note that the page headers are aligned towards the outside of the page (right on the right-hand page, left on the left-hand page) and they contain the section title on the right and the chapter title on the left respectively, in **SIMLLCAPS**. The footers contain only page numbers on the exterior of the page, aligned right or left depending on the page. The lines used to delimit the headers and footers from the rest of the page are **0.4pt** thick, and are as long as the text.

Table A.2: Font types, faces and sizes to be used.

Element	Face	Size	\LaTeX size
Ch. label	serif, bold	24.88pt	<code>\huge</code>
Chapter	serif, bold	24.88pt	<code>\Huge</code>
Section	sans-serif, bold	20.74pt	<code>\LARGE</code>
Subsection	sans-serif, bold	17.28pt	<code>\Large</code>
Subsubsection	sans-serif, bold	14.4pt	<code>\large</code>
Body	serif	12pt	<code>\normalsize</code>
HEADER	SERIF, SMALLCAPS	10pt	
Footer (page numbers)	serif, regular	12pt	
Figure label	serif, bold	12pt	
Figure caption	serif, regular	12pt	
In figure	sans-serif	any	
Table label	serif, bold	12pt	
Table caption and text	serif, regular	12pt	
Listings	mono-space	\leq 12pt	

A.2.2 Chapters, Sections, Paragraphs

Chapter, section, subsection, etc. names are all left aligned, and numbered as in this document.

Chapters always start on the right-hand page, with the label and title separated from the rest of the text by a **0.4pt** thick line.

Paragraphs are justified (left and right), using single line spacing. Note that the first paragraph of a chapter, section, etc. is not indented, while the following are indented.

A.2.3 Tables

Table captions should be located above the table, justified, and spaced 2.0cm from left and right (important for very long captions). Tables should be numbered, but the numbering is up to you, and could be, for instance:

- **Table XY** where X is the chapter number and Y is the table number within that chapter. (This is the default in \LaTeX . More on \LaTeX can be found on-line, including whole books, such as [?].) or
- **Table Y** where Y is the table number within the whole report

As a recommendation, use regular paragraph text in the tables, bold headings and avoid vertical lines (see Table A.2).

A.2.4 Figures

Figure labels, numbering, and captions should be formed similarly to tables. As a recommendation, use vector graphics in figures (Figure A.1), rather than bitmaps (Figure A.2). Text within figures usually looks better with sans-serif fonts.

This is vector graphics



Figure A.1: A PDF vector graphics figure. Notice the numbering and placement of the caption. The caption text is indented 2.0cm from both left and right text margin.

This is raster graphics



Figure A.2: A JPEG bitmap figure. Notice the bad quality of such an image when scaling it. Sometimes bitmap images are unavoidable, such as for screen dumps.

For those interested in delving deeper into the design of graphical information display, please refer to books such as [?, ?].

A.3 Mathematical Formulae and Equations

You are free to use in-text equations and formulae, usually in *italic serif* font. For instance: $S = \sum_i a_i$. We recommend using numbered equations when you do need to refer to the

specific equations:

$$E = \int_0^\delta P(t)dt \quad \longleftrightarrow \quad E = mc^2 \quad (\text{A.1})$$

The numbering system for equations should be similar to that used for tables and figures.

A.4 References

Your references should be gathered in a **References** section, located at the end of the document (before **Appendices**). We recommend using number style references, ordered as appearing in the document or alphabetically. Have a look at the references in this template in order to figure out the style, fonts and fields. Web references are acceptable (with restraint) as long as you specify the date you accessed the given link [?, ?]. You may of course use URLs directly in the document, using mono-space font, i.e. <http://cs.lth.se/>.

Make sure you add references as close to the claim as possible [?], as shown, not at the end of a whole paragraph. Notice also that there is a space before the reference; best is to use `\cite{ref}`, to allow for unbreakable spaces. References should not be used after the period marking the end of sentence. Using the reference as follows (end of paragraph, after period) is *strongly discouraged*, since it says nothing about which specific claim you provide the reference for. [?]

A.5 Colours

As a general rule, all theses are printed in black-and-white, with the exception of selected parts in selected theses that need to display colour images essential to describing the thesis outcome (*computer graphics*, for instance).

A strong requirement is for using **black text on white background** in your document's main text. Otherwise we do encourage using colours in your figures, or other elements (i.e. the colour marking internal and external references) that would make the document more readable on screen. You may also emphasize table rows, columns, cells, or headers using white text on black background, or black text on light grey background.

Note that the document should look good in black-and-white print. Colours are often rendered using monochrome textures in print, which makes them look different from on screen versions. This means that you should choose your colours wisely, and even opt for black-and-white textures when the distinction between colours is hard to make in print. The best way to check how your document looks, is to print out a copy yourself.

The L^AT_EX class defines also a few *LTH* standard colours, which you could use in your document for various elements, to adhere to the standard university profile.

These are: `LTHblue`, `LTHbronze`, `LTHgreen`, `LTHpink`, `LTHcyan`, `LTHgrey`.

Appendix B

Language

You are strongly encouraged to write your report in English, for two reasons. First, it will improve your use of English language. Second, it will increase visibility for you, the author, as well as for the Department of Computer Science, and for your host company (if any).

However, note that your examiner (and supervisors) are not there to provide you with extensive language feedback. We recommend that you check the language used in your report in several ways:

Reference books dedicated to language issues can be very useful. [?]

Spelling and grammar checkers which are usually available in the commonly used text editing environments.

Colleagues and friends willing to provide feedback your writing.

Studieverkstaden is a university level workshop, that can help you with language related problems (see Studieverkstaden's web page).

Websites useful for detecting language errors or strange expressions, such as

- <http://translate.google.com>
- <http://www.gingersoftware.com/grammarcheck/>

B.1 Style Elements

Next, we will just give some rough guidelines for good style in a report written in English. Your supervisor and examiner as well as the aforementioned **Studieverkstad** might have a different take on these, so we recommend you follow their advice whenever in doubt. If you want a reference to a short style guide, have a look at [?].

Widows and Orphans

Avoid *widows* and *orphans*, namely words or short lines at the beginning or end of a paragraph, which are left dangling at the top or bottom of a column, separated from the rest of the paragraph.

Footnotes

We strongly recommend you avoid footnotes. To quote from [?], *Footnotes are frequently misused by containing information which should either be placed in the text or excluded altogether. They should be avoided as a general rule and are acceptable only in exceptional cases when incorporation of their content in the text [is] not possible.*

Active vs. Passive Voice

Generally active voice (*I ate this apple.*) is easier to understand than passive voice (*This apple has been eaten (by me).*) In passive voice sentences the actor carrying out the action is often forgotten, which makes the reader wonder who actually performed the action. In a report is important to be clear about who carried out the work. Therefore we recommend to use active voice, and preferably the plural form *we* instead of *I* (even in single author reports).

Long and Short Sentences

A nice brief list of sentence problems and solutions is given in [?]. Using choppy sentences (too short) is a common problem of many students. The opposite, using too long sentences, occurs less often, in our experience.

Subject-Predicate Agreement

A common problem of native Swedish speakers is getting the subject-predicate (verb) agreement right in sentences. Note that a verb must agree in person and number with its subject. As a rough tip, if you have subject ending in *s* (plural), the predicate should not, and the other way around. Hence, *only one s*. Examples follow:

incorrect He have to take this road.

correct He has to take this road.

incorrect These words forms a sentence.

correct These words form a sentence.

In more complex sentences, getting the agreement right is trickier. A brief guide is given in the *20 Rules of Subject Verb Agreement* [?].

Appendix C

Structure

It is a good idea to discuss the structure of the report with your supervisor rather early in your writing. Given next is a generic structure that is a starting point, but by no means the absolute standard. Your supervisor should provide a better structure for the specific field you are writing your thesis in. Note also that the naming of the chapters is not compulsory, but may be a helpful guideline.

Introduction should give the background of your work. Important parts to cover:

- Give the context of your work, have a short introduction to the area.
- Define the problem you are solving (or trying to solve).
- Specify your contributions. What does this particular work/report bring to the research area or to the body of knowledge? How is the work divided between the co-authors? (This part is essential to pinpoint individual work. For theses with two authors, it is compulsory to identify which author has contributed with which part, both with respect to the work and the report.)
- Describe related work (literature study). Besides listing other work in the area, mention how it is related or relevant to your work. The tradition in some research area is to place this part at the end of the report (check with your supervisor).

Approach should contain a description of your solution(s), with all the theoretical background needed. On occasion this is replaced by a subset or all of the following:

- **Method:** describe how you go about solving the problem you defined. Also how do you show/prove that your solution actually works, and how well does it work.
- **Theory:** should contain the theoretical background needed to understand your work, if necessary.
- **Implementation:** if your work involved building an artefact/implementation, give the details here. Note, that this should not, as a rule, be a chronological

description of your efforts, but a view of the result. There is a place for insights and lamentation later on in the report, in the Discussion section.

Evaluation is the part where you present the finds. Depending on the area this part contains a subset or all of the following:

- **Experimental Setup** should describe the details of the method used to evaluate your solution(s)/approach. Sometimes this is already addressed in the **Method**, sometimes this part replaces **Method**.
- **Results** contains the data (as tables, graphs) obtained via experiments (benchmarking, polls, interviews). Here you should also describe the individual tables or graphs in text, pointing out interesting outliers and trends.
- **Discussion** allows for a longer discussion and interpretation of the results from the evaluation, including extrapolations and/or expected impact. Focus here on a broader view of the results, talking about the relation between the different finds.¹ This might also be a good place to describe your positive and negative experiences related to the work you carried out.

Occasionally these sections are intermingled, if this allows for a better presentation of your work. However, try to distinguish between measurements or hard data (results) and extrapolations, interpretations, or speculations (discussion).

Conclusions should summarize your findings and possible improvements or recommendations.

Bibliography is a must in a scientific report. L^AT_EX and **bibtex** offer great support for handling references and automatically generating bibliographies.

Appendices should contain lengthy details of the experimental setup, mathematical proofs, code download information, and shorter code snippets. Avoid longer code listings. Source code should rather be made available for download on a website or on-line repository of your choosing.

¹Bad practice is to display graphs in Results and then describe them textually one by one in here. No! Both sections should have some discussion, but one targets individual finds and the other tries to bridge between these adopting a more overarching viewpoint.

EXAMENSARBETE Application Specific Instruction-set Processor Using a Parametrizable multi-SIMD
Synthesizeable Model Supporting Design Space Exploration
STUDENT Magnus Hultin
HANLEDARE Flavius Gruian (LTH)
EXAMINATOR Krzysztof Kuchcinski (LTH)

Parametrisk processor modell för design utforskning

POPULÄRVETENSKAPLIG SAMMANFATTNING **Magnus Hultin**

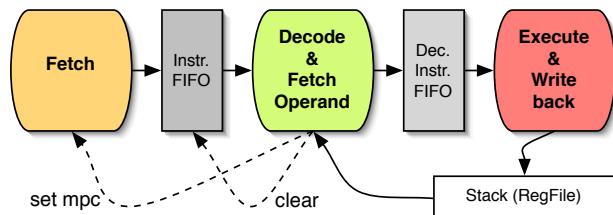
Applikations-specifika processorer är allt mer vanligt för få ut rätt prestanda med så lite resurser som möjligt. Detta arbete har en parametrisk modell för att kunna testa hur mycket resurser som behövs för en specifik applikation.

För att öka prestandan i dagens processorer finns det vektorenheter och flera kärnor i processorer. Vektorenheten finns till för att kunna utföra en operation på en mängd data samtidigt och flera kärnor gör att man kan utföra fler instruktioner samtidigt. Ofta är processorerna designade för att kunna stödja en mängd olika datorprogram. Detta resulterar i att det blir kompromisser som kan påverka prestandan för vissa program och vara överflödig för andra. I t.ex. videokameror, mobiltelefoner, medicinsk utrustning, digital kameror och annan inbyggd elektronik, kan man istället använda en processor som saknar vissa funktioner men som istället är mer energieffektiv. Man kan jämföra det med att frakta ett paket med en stor lastbil istället för att använda en mindre bil där samma paketet också skulle få plats.

I mitt examensarbete har jag skrivit en modell som kan användas för att snabbt designa en processor enligt vissa parametrar. Dessa parametrar väljs utifrån vilket eller vilka program man tänkt köra på den. Vissa program kan t.ex. lättare använda flera kärnor och vissa program kan använda korta eller längre vektorenheter för dess data.

Modellen testades med olika multimedia program. Den mest beräkningsintensiva och mest up-

prende delen av programmen användes. Dessa kallas för kärnor av programmen. Kärnorna som användes var ifrån MPEG och JPEG, som används för bildkomprimering och videokomprimering.



Resultatet visar att det finns en prestanda vinst jämfört med generella processorer men att detta också ökar resurserna som behövs. Detta trots att den generella processorn har nästan dubbelt så hög klockfrekvens än dem applikations-specifika processorer. Resultatet visar också att schemaläggning av instruktionerna i programmen spelar en stor roll för att kunna utnyttja resurserna som finns tillgängliga och därmed öka prestandan. Med den schemaläggningen som utnyttjade resurserna bäst var prestandan minst 79% bättre än den generella processorn.